

1. Create a DataFrame using the following data `data = {'Name': ['Alice', 'Bob', 'Charlie', 'David'], 'Age': [24, 27, 22, 32], 'Salary': [50000, 60000, 45000, 80000]}` Write a program to: a) Create a new column, Tax, which is 10% of the Salary. b) Create another column, Net Salary, as `Salary - Tax`.

```
import pandas as pd

# Creating a DataFrame
data = {
    'Name': ['Alice', 'Bob', 'Charlie'],
    'Age': [25, 30, 35],
    'City': ['New York', 'Los Angeles', 'Chicago']
}

df = pd.DataFrame(data)
print(df)
```

```
↩
```

	Name	Age	City
0	Alice	25	New York
1	Bob	30	Los Angeles
2	Charlie	35	Chicago

```
print(df['Name'])
```

```
↩
```

0	Alice
1	Bob
2	Charlie

Name: Name, dtype: object

```
print(df.iloc[0]) # Access by position
print(df.loc[0]) # Access by label (index)
```

```
↩
```

Name	Alice
Age	25
City	New York

Name: 0, dtype: object

Name	Alice
Age	25
City	New York

Name: 0, dtype: object

2. Given the following DataFrame `data = {'Sex': ['male', 'female', 'female', 'male', 'female'], 'Class': ['First', 'Second', 'Third', 'First', 'Second'], 'Fare': [100, 50, 20, 120, 60]}` `df = pd.DataFrame(data)` Write a program to calculate the mean ticket fare price for each combination of Sex and Class.

```
import pandas as pd

# Sample DataFrame
data = {'Name': ['Alice', 'Bob', 'Charlie'],
        'Age': [24, 27, 22],
        'City': ['New York', 'Los Angeles', 'Chicago']}
df = pd.DataFrame(data, index=['A', 'B', 'C'])
print(df, "\n")
# Using loc (label-based)
print(df.loc['A']) # Access row with index label 'A'
print("\n")
print(df.loc['A', 'Age']) # Access 'Age' column for row 'A'
print("\n")
print(df.loc['A':'B', 'Name']) # Slice rows 'A' to 'B' and 'Name' column
```

```
↩
```

	Name	Age	City
A	Alice	24	New York
B	Bob	27	Los Angeles
C	Charlie	22	Chicago

Name	Alice
Age	24
City	New York

Name: A, dtype: object

24

```
A    Alice
B     Bob
Name: Name, dtype: object
```

```
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David'],
    'Age': [24, 27, 22, 32],
    'Salary': [50000, 60000, 45000, 80000]
}
df = pd.DataFrame(data);
print(df)
```

```
↔
```

	Name	Age	Salary
0	Alice	24	50000
1	Bob	27	60000
2	Charlie	22	45000
3	David	32	80000

```
df['Tax'] = df['Salary']*0.1
print(df)
```

```
↔
```

	Name	Age	Salary	Tax
0	Alice	24	50000	5000.0
1	Bob	27	60000	6000.0
2	Charlie	22	45000	4500.0
3	David	32	80000	8000.0

```
df['net_salary'] = df['Salary']-df['Tax']
print(df)
```

```
↔
```

	Name	Age	Salary	Tax	net_salary
0	Alice	24	50000	5000.0	45000.0
1	Bob	27	60000	6000.0	54000.0
2	Charlie	22	45000	4500.0	40500.0
3	David	32	80000	8000.0	72000.0

Double-click (or enter) to edit

```
data = {
    'Sex': ['male', 'female', 'female', 'male', 'female'],
    'Class': ['First', 'Second', 'Third', 'First', 'Second'],
    'Fare': [100, 50, 20, 120, 60]
}
df = pd.DataFrame(data)
print(df)
```

```
↔
```

	Sex	Class	Fare
0	male	First	100
1	female	Second	50
2	female	Third	20
3	male	First	120
4	female	Second	60

```
male_fare = df.loc[df['Sex']=='male','Fare']
female_fare = df.loc[df['Sex']=='female','Fare']
male_fare_mean = np.mean(male_fare);
female_fare_mean = np.mean(female_fare)
print(male_fare)
print(female_fare)
print(male_fare_mean)
print(round(female_fare_mean,2))
```

```
↔
```

0	100
3	120

Name: Fare, dtype: int64

1	50
2	20
4	60

Name: Fare, dtype: int64

110.0

43.33

3. Using the same DataFrame as Question 2, write a program to count the number of passengers in each Class.

```
class_count = df['Class'].value_counts()
print(class_count)
```

```
↗ Class
First      2
Second     2
Third      1
Name: count, dtype: int64
```

4. Given the following DataFrame data = { 'Student': ['Alice', 'Bob', 'Charlie'], 'Math': [85, 90, 95], 'Science': [88, 92, 96] } df = pd.DataFrame(data)

Write a program to reshape the DataFrame from wide format to long format such that each row corresponds to a Student, a Subject, and their respective Score.

```
import pandas as pd

# Original DataFrame
data = {
    'Student': ['Alice', 'Bob', 'Charlie'],
    'Math': [85, 90, 95],
    'Science': [88, 92, 96]
}
df = pd.DataFrame(data)

# Reshape the DataFrame to long format
df_long = pd.melt(df, id_vars=['Student'], var_name='Subject', value_name='Score')

# Display the reshaped DataFrame
print(df_long)
```

```
↗
```

	Student	Subject	Score
0	Alice	Math	85
1	Bob	Math	90
2	Charlie	Math	95
3	Alice	Science	88
4	Bob	Science	92
5	Charlie	Science	96

5. Using the reshaped DataFrame from Question 4, write a program to convert it back to wide format with Math and Science as separate columns.

```
import pandas as pd

# Reshaped (long format) DataFrame
data_long = {
    'Student': ['Alice', 'Bob', 'Charlie', 'Alice', 'Bob', 'Charlie'],
    'Subject': ['Math', 'Math', 'Math', 'Science', 'Science', 'Science'],
    'Score': [85, 90, 95, 88, 92, 96]
}
df_long = pd.DataFrame(data_long)

# Convert back to wide format
df_wide = df_long.pivot(index='Student', columns='Subject', values='Score').reset_index()

# Display the wide format DataFrame
print(df_wide)
```

```
↗
```

	Student	Math	Science
0	Alice	85	88
1	Bob	90	92
2	Charlie	95	96

6. Given the following two DataFrames data1 = {'ID': [1, 2, 3], 'Name': ['Alice', 'Bob', 'Charlie']} data2 = {'ID': [2, 3, 4], 'Score': [85, 90, 95]} df1 = pd.DataFrame(data1) df2 = pd.DataFrame(data2) Write a program to: a) Perform an inner join on ID. b) Perform a left join on ID.

```
# DataFrames
data1 = {'ID': [1, 2, 3], 'Name': ['Alice', 'Bob', 'Charlie']}
data2 = {'ID': [2, 3, 4], 'Score': [85, 90, 95]}
df1 = pd.DataFrame(data1)
df2 = pd.DataFrame(data2)

# a) Perform an inner join on ID
inner_join_result = pd.merge(df1, df2, on='ID', how='inner')

# b) Perform a left join on ID
left_join_result = pd.merge(df1, df2, on='ID', how='left')

# Display the results
print("Inner Join Result:")
print(inner_join_result)

print("\nLeft Join Result:")
print(left_join_result)
```

```
↔ Inner Join Result:
   ID  Name  Score
0    2   Bob    85
1    3  Charlie   90

Left Join Result:
   ID  Name  Score
0    1  Alice   NaN
1    2   Bob   85.0
2    3  Charlie  90.0
```

7. Given the following two DataFrames data1 = {'ID': [1, 2, 3], 'Name': ['Alice', 'Bob', 'Charlie']} data2 = {'ID': [2, 3, 4], 'Score': [85, 90, 95]} df1 = pd.DataFrame(data1) df2 = pd.DataFrame(data2) Write a program to: a) Perform an inner join on ID. b) Perform a left join on ID.

```
# Define DataFrames
data1 = {'ID': [1, 2, 3], 'Name': ['Alice', 'Bob', 'Charlie']}
data2 = {'ID': [2, 3, 4], 'Score': [85, 90, 95]}
df1 = pd.DataFrame(data1)
df2 = pd.DataFrame(data2)

# a) Perform an inner join on ID
inner_join_result = pd.merge(df1, df2, on='ID', how='inner')

# b) Perform a left join on ID
left_join_result = pd.merge(df1, df2, on='ID', how='left')

# Display the results
print("Inner Join Result:")
print(inner_join_result)

print("\nLeft Join Result:")
print(left_join_result)
```

```
↔ Inner Join Result:
   ID  Name  Score
0    2   Bob    85
1    3  Charlie   90

Left Join Result:
   ID  Name  Score
0    1  Alice   NaN
1    2   Bob   85.0
2    3  Charlie  90.0
```

8. Using the following DataFrame: data = {'Department': ['HR', 'IT', 'Finance', 'HR', 'Finance', 'IT'], 'Employee': ['Alice', 'Bob', 'Charlie', 'David', 'Eve', 'Frank'], 'Salary': [50000, 60000, 55000, 52000, 59000, 61000]} df = pd.DataFrame(data) Write a program to: a) Calculate the total salary for each department. b) Calculate the average salary for each department

```
# Define the DataFrame
data = {
    'Department': ['HR', 'IT', 'Finance', 'HR', 'Finance', 'IT'],
    'Employee': ['Alice', 'Bob', 'Charlie', 'David', 'Eve', 'Frank'],
```

```

'Salary': [50000, 60000, 55000, 52000, 59000, 61000]
}
df = pd.DataFrame(data)

# a) Calculate the total salary for each department
total_salary = df.groupby('Department')['Salary'].sum()

# b) Calculate the average salary for each department
average_salary = df.groupby('Department')['Salary'].mean()

# Display the results
print("Total Salary for Each Department:")
print(total_salary)

print("\nAverage Salary for Each Department:")
print(average_salary)

```

```

↗ Total Salary for Each Department:
Department
Finance    114000
HR          102000
IT          121000
Name: Salary, dtype: int64

Average Salary for Each Department:
Department
Finance     57000.0
HR          51000.0
IT          60500.0
Name: Salary, dtype: float64

```

9. Given the following DataFrame: data = { 'Name': ['Alice', 'Bob', 'Charlie', 'David'], 'Age': [24, np.nan, 22, 32], 'Salary': [50000, 60000, np.nan, 80000] } df = pd.DataFrame(data) Write a program to: a) Replace the missing values in the Age column with the mean age. b) Drop rows where the Salary column has missing value

```

# Define the DataFrame
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David'],
    'Age': [24, np.nan, 22, 32],
    'Salary': [50000, 60000, np.nan, 80000]
}
df = pd.DataFrame(data)

# a) Replace missing values in the Age column with the mean age
mean_age = df['Age'].mean()
df['Age'].fillna(mean_age, inplace=True)

# b) Drop rows where the Salary column has missing values
df = df.dropna(subset=['Salary'])

# Display the results
print("DataFrame after replacing missing Age values and dropping Salary rows:")
print(df)

```

10. Using the following DataFrames data = { 'Month': ['Jan', 'Feb', 'Mar', 'Apr', 'May'], 'Sales': [300, 400, 500, 600, 700] }

df = pd.DataFrame(data) Write a program to create: a) A line plot for Month vs. Sales. b) A bar plot for the same data.

```

import matplotlib.pyplot as plt

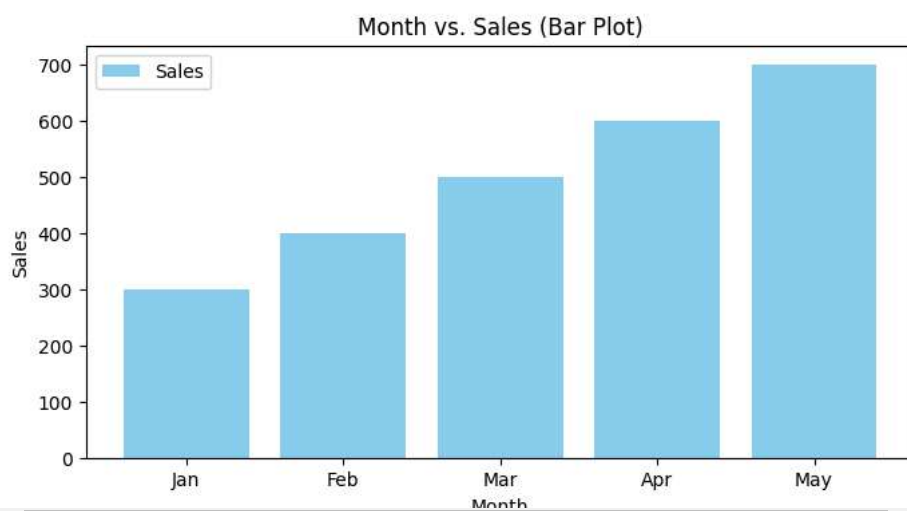

# Define the DataFrame
data = { 'Month': ['Jan', 'Feb', 'Mar', 'Apr', 'May'], 'Sales': [300, 400, 500, 600, 700] }
df = pd.DataFrame(data)

# a) Line plot for Month vs. Sales
plt.figure(figsize=(8, 4))
plt.plot(df['Month'], df['Sales'], marker='o', label='Sales')
plt.title('Month vs. Sales (Line Plot)')
plt.xlabel('Month')
plt.ylabel('Sales')
plt.grid(True)

```

```
plt.legend()
plt.show()
```

```
# b) Bar plot for Month vs. Sales
plt.figure(figsize=(8, 4))
plt.bar(df['Month'], df['Sales'], color='skyblue', label='Sales')
plt.title('Month vs. Sales (Bar Plot)')
plt.xlabel('Month')
plt.ylabel('Sales')
plt.legend()
plt.show()
```



11. Create a DataFrame with a Date column containing the following dates ['2023-01-01', '2023-01-02', '2023-01-03', '2023-01-04', '2023-01-05']
Add a column Sales with random integers between 100 and 500. Write a program to: a) Convert the Date column to a datetime object. b)
Filter the rows where Sales are greater than 300.

```
import numpy as np
# Create the DataFrame
data = {
    'Date': ['2023-01-01', '2023-01-02', '2023-01-03', '2023-01-04', '2023-01-05'],
    'Sales': np.random.randint(100, 500, size=5)
}
df = pd.DataFrame(data)

# a) Convert the Date column to a datetime object
df['Date'] = pd.to_datetime(df['Date'])

# b) Filter rows where Sales are greater than 300
filtered_df = df[df['Sales'] > 300]

# Display the results
```

```
print("Original DataFrame:")
print(df)

print("\nFiltered DataFrame (Sales > 300):")
print(filtered_df)
```

```
Original DataFrame:
   Date  Sales
0 2023-01-01   280
1 2023-01-02   456
2 2023-01-03   157
3 2023-01-04   215
4 2023-01-05   118

Filtered DataFrame (Sales > 300):
   Date  Sales
1 2023-01-02   456
```

12. Create a series of 4 capital cities where the index is the name of corresponding country.

```
import pandas as pd

# Create the Series
capitals = pd.Series(
    data=['Washington, D.C.', 'London', 'Paris', 'Tokyo'],
    index=['USA', 'UK', 'France', 'Japan']
)

# Display the Series
print(capitals)
```

```
USA      Washington, D.C.
UK         London
France      Paris
Japan       Tokyo
dtype: object
```

13. Create a dataframe of (at least 4) countries, with 2 variables: population and capital. Country name should be the index.

```
# Create the DataFrame
data = {
    'Population': [331002651, 67886011, 65273511, 126476461], # Population values
    'Capital': ['Washington, D.C.', 'London', 'Paris', 'Tokyo'] # Capitals
}
countries = pd.DataFrame(data, index=['USA', 'UK', 'France', 'Japan'])

# Display the DataFrame
print(countries)
```

```
Population      Capital
USA      331002651  Washington, D.C.
UK         67886011      London
France     65273511      Paris
Japan      126476461      Tokyo
```

14. How many columns are printed? How many variables does the dataframe contain? Dataset: titanic.csv

```
# Load the Titanic dataset
titanic_df = pd.read_csv('path_to_titanic.csv') # Provide the correct file path

# Get the number of columns and variables
num_columns = len(titanic_df.columns)
num_variables = titanic_df.shape[1]

print(f"Number of columns: {num_columns}")
print(f"Number of variables: {num_variables}")
```

15. Create a matrix of data, and create a data frame from it using pd.DataFrame. Specify index (row names) and columns (variable names). Include at least 3 cities and 3 variables (e.g. population in millions, size in km2, and population density people per km2).

```
# Create a matrix of data
data = [
    [8.4, 789, 10600], # City 1 (New York): population in millions, size in km^2, population density
    [3.9, 607, 6400], # City 2 (London)
    [14.5, 174, 83000] # City 3 (Tokyo)
]

# Create the DataFrame
cities = ['New York', 'London', 'Tokyo']
columns = ['Population (millions)', 'Size (km^2)', 'Population Density (people per km^2)']
df = pd.DataFrame(data, index=cities, columns=columns)

# Display the DataFrame
print(df)
```

```
↗
Population (millions)  Size (km^2)  \
New York              8.4          789
London               3.9          607
Tokyo              14.5          174

Population Density (people per km^2)
New York              10600
London               6400
Tokyo              83000
```

16. Take your own city matrix and city data frame. From both of these extracts: a) population density (for all cities) b) data for the third city. For the data frame do it in two ways: using index, and using row number!

```
# Create a matrix of data for 3 cities
data = [
    [5.2, 400, 13000], # City A: population in millions, size in km^2, population density
    [3.1, 500, 6200], # City B
    [9.8, 600, 16333] # City C
]

# Create the DataFrame
cities = ['City A', 'City B', 'City C']
columns = ['Population (millions)', 'Size (km^2)', 'Population Density (people per km^2)']
df = pd.DataFrame(data, index=cities, columns=columns)

# a) Population density for all cities
population_density = df['Population Density (people per km^2)']

# b) Data for the third city (City C)
# 1) Using index (by city name)
city_c_data_by_index = df.loc['City C']

# 2) Using row number (3rd row in the DataFrame)
city_c_data_by_row_number = df.iloc[2]

# Display the results
print("Population Density for all cities:")
print(population_density)

print("\nData for the third city (City C) using index:")
print(city_c_data_by_index)

print("\nData for the third city (City C) using row number:")
print(city_c_data_by_row_number)
```

```
↗ Population Density for all cities:
City A    13000
City B     6200
City C    16333
Name: Population Density (people per km^2), dtype: int64

Data for the third city (City C) using index:
Population (millions)    9.8
Size (km^2)             600.0
Population Density (people per km^2)    16333.0
Name: City C, dtype: float64

Data for the third city (City C) using row number:
Population (millions)    9.8
```



```
Size (km²)          600.0
Population Density (people per km²)  16333.0
Name: City C, dtype: float64
```